

README.md

2024-02-14

```
knitr::opts_chunk$set(echo = TRUE)

relu <- function(x) {
  # Rectified linear unit activation function
  # Replaces negative values with zero
  x[x < 0] <- 0
  x
}

softmax <- function(x) {
  # Softmax activation function
  # Normalizes the input to a probability distribution
  exp_x <- exp(x)
  row_sums <- apply(exp_x, 1, sum)
  exp_x / row_sums
}

one_hot_encode <- function(x) {
  # One-hot encoding function
  # Converts a vector of labels into a matrix of binary indicators
  n <- length(x)
  k <- max(x)
  mat <- matrix(0, nrow = n, ncol = k)
  mat[cbind(1:n, x)] <- 1
  mat
}

# Initialize the parameters with random values
initialize <- function(n_hidden, n_features = 2, n_class = 2) {
  return(list(
    W1 = matrix(rnorm(n_features * n_hidden), nrow = n_features, ncol = n_hidden),
    b1 = rnorm(n_hidden),
    W2 = matrix(rnorm(n_hidden * n_class), nrow = n_hidden, ncol = n_class),
    b2 = rnorm(n_class)
  ))
}

# forwardPropagate pass of the neural network
forwardPropagate <- function(input_data, params) {
  W1 <- params$W1
  W2 <- params$W2
  b1 <- params$b1
  b2 <- params$b2
```

```

    a1 <- (input_data %*% W1) + b1
    params$z1 <- relu(a1)
    a2 <- (params$z1 %*% W2) + b2
    params$z2 <- softmax(a2)

    return(list(params$z1, params$z2))
}

# Fit the neural network to the input data and labels
fit = function(input_data, label, batch_size, iter_num, params) {
  for (epoch in 1:iter_num) { #Iterate through the requested number of iterations
    p <- sample(1:length(label)) # Use sample to generate a random int
    input_data <- input_data[p, ] # select the row
    label <- label[p] #select the label
    for (i in seq(1, length(label), by = batch_size)) {
      batch_data <- input_data[i:(i + batch_size - 1), ]
      batch_label <- label[i:(i + batch_size - 1)]
      params <- sgd(batch_data, batch_label, params) # We update the params
    }
  }
  return(params)
}

# Stochastic gradient descent update of the parameters
sgd <- function(data, label, params, alpha = 1e-4) {
  grad <- backward(data, label, params)
  for (layer in names(grad)) {
    params[[layer]] <- params[[layer]] + (alpha * grad[[layer]])
  }
  return(params)
}

# Backward pass of the neural network
backward <- function(data, label, params) {
  W1 <- params$W1
  W2 <- params$W2
  b1 <- params$b1
  b2 <- params$b2
  z1 <- forwardPropagate(data, params)[[1]]
  z2 <- forwardPropagate(data, params)[[2]]

  #label <- one_hot_encode(label)
  db2_temp <- label - z2
  db2 <- colSums(db2_temp)
  dW2 <- t(z1) %*% db2_temp
  db1_temp <- db2_temp %*% t(W2)
  db1_temp[z1 <= 0] <- 0
  db1 <- colSums(db1_temp)
  dW1 <- t(data) %*% db1_temp

  return(list(W1 = dW1, b1 = db1, W2 = dW2, b2 = db2))
}

```

```

# Test the accuracy of the neural network on the test data and labels
test = function(train_data, train_label, test_data, test_label, batch_size, iter_num, params) {
  pred_label <- forwardPropagate(test_data, params)[[2]]
  pred_label <- apply(pred_label, 1, which.max)
  acc <- mean(pred_label == test_label)
  return(acc)
}

set.seed(123)
n_train <- 100
n_test <- 20
n_features <- 2
n_class <- 2
train_data <- matrix(rnorm(n_train * n_features), nrow = n_train, ncol = n_features)
train_label <- sample(1:n_class, n_train, replace = TRUE)
test_data <- matrix(rnorm(n_test * n_features), nrow = n_test, ncol = n_features)
test_label <- sample(1:n_class, n_test, replace = TRUE)

# Instantiate the NeuralNet class with 10 hidden units
nn <- initialize(n_hidden = 10)

# Train the neural network with batch size of 10 and 50 iterations
nn <- fit(train_data, train_label, batch_size = 10, iter_num = 1, nn)

# Test the accuracy of the neural network on the test data
acc <- test(train_data, train_label, test_data, test_label, batch_size = 10, iter_num = 50, nn)
cat("Accuracy:", acc, "\n")

```

```
## Accuracy: 0.5
```

```

knitr::opts_chunk$set(echo = TRUE)

# The neural network so to speak is just weights and biases
set.seed(123)

relu <- function(x) {
  # Rectified linear unit activation function
  # Replaces negative values with zero
  x[x < 0] <- 0
  x
}

softmax <- function(x) {
  # Softmax activation function
  # Normalizes the input to a probability distribution
  exp_x <- exp(x)
  row_sums <- apply(exp_x, 1, sum)
  exp_x / row_sums
}

```

```

}

one_hot_encode <- function(x) {
  # One-hot encoding function
  # Converts a vector of labels into a matrix of binary indicators
  n <- length(x)
  #k <- max(x)
  k <- length(unique(x))
  mat <- matrix(0, nrow = n, ncol = k)
  mat[cbind(1:n, x)] <- 1
  mat
}

one_hot_encode <- function(x) {
  # One-hot encoding function
  # Converts a vector of string labels into a matrix of binary indicators
  n <- length(x)
  k <- length(unique(x))
  labels <- unique(x)
  mat <- matrix(0, nrow = n, ncol = k)
  for (i in 1:n) {
    j <- which(labels == x[i])
    mat[i, j] <- 1
  }
  mat
}

string_to_int <- function(x) {
  unique_strings <- unique(x)
  int_values <- seq_along(unique_strings)
  x_int <- int_values[match(x, unique_strings)]
  return(x_int)
}

n_train <- 100
n_test <- 20
n_features <- 4
n_class <- 3

## generate a random ordering
set.seed(1) ## make reproducible here, but not if generating many random samples
rand <- sample(nrow(iris))
rand

```

```

##      [1]  68 129  43  14  51  85  21 106  74   7  73  79  37 105 110  34 143 126
##     [19]  89  33  84  70 142  42  38 111  20  28 124  44  87 149  40 121  25 119
##     [37]  39 146 127   6  24  32 147   2  45  18  22  78 102  65 115 120 100  75
##     [55]  81  13 118 132  48  93  23 130  29  95 104 123  92 131 134 144  31  17
##     [73] 140  91  64  60 113 135  10   1 148  59  26  15  58  88 136 112  77  53
##     [91]  12 114  76  61 145  86  94  83  19 150  35  98  71 101 108  55 125  56

```

```
## [109] 41 138 3 82 50 141 133 27 63 122 116 66 52 96 16 30 4 107
## [127] 49 62 97 103 9 99 109 137 54 90 139 11 80 69 36 8 67 46
## [145] 128 47 117 5 57 72
```

```
d <- iris[rand,]
train_data <- as.matrix(d[0:75,0:4])

train_label <- one_hot_encode(as.matrix(d[0:75,5:5]))

test_data <- as.matrix(d[75:150,0:4])

test_label <- one_hot_encode(as.matrix(d[75:150,5:5]))

n_hidden = 10

params = list(
  W1 = matrix(rnorm(n_features * n_hidden), nrow = n_features, ncol = n_hidden),
  b1 = rnorm(n_hidden),
  W2 = matrix(rnorm(n_hidden * n_class), nrow = n_hidden, ncol = n_class),
  b2 = rnorm(n_class)
)

iter_num = 10
batch_size = 10
alpha = 1e-4

# Backward pass of the neural network
backward <- function(data, label, params) {
  W1 <- params$W1
  W2 <- params$W2
  b1 <- params$b1
  b2 <- params$b2
  z1 <- forwardPropogate(data, params)[[1]]
  z2 <- forwardPropogate(data, params)[[2]]

  db2_temp <- label - z2
  db2 <- colSums(db2_temp)
  dW2 <- t(z1) %*% db2_temp
  db1_temp <- db2_temp %*% t(W2)
  db1_temp[z1 <= 0] <- 0
  db1 <- colSums(db1_temp)
  dW1 <- t(data) %*% db1_temp

  return(list(W1 = dW1, b1 = db1, W2 = dW2, b2 = db2))
}

forwardPropogate <- function( input_data,params) {
  W1 <- params$W1
  W2 <- params$W2
  b1 <- params$b1
  b2 <- params$b2

  a1 <- (input_data %*% W1) + b1
```

```

    params$z1 <- relu(a1)
    a2 <- (params$z1 %*% W2) + b2
    params$z2 <- softmax(a2)

    return(list(params$z1, params$z2))
  }

train_dim <- dim(train_label)[1]

s <- seq(1, train_dim, by = batch_size)

y <- pmin(s+batch_size, train_dim)
l <- seq(1, length(y))

for (epoch in 1:iter_num) { #Iterate through the requested number of iterations
  p <- sample(1:train_dim) # Use sample to generate a random int
  input_data <- train_data[p, ] # select the row
  label <- train_label[p] #select the label
  for (k in l) {
    i <- s[k]
    j <- y[k]

    batch_data <- train_data[i:j, ]

    batch_label <- train_label[i:j]

    # sgd <- function(data, label, params, alpha = 1e-4) {

    grad <- backward(batch_data, batch_label, params)

    for (layer in names(grad)) {
      params[[layer]] <- params[[layer]] + (alpha * grad[[layer]])
    }

    # params <- sgd(batch_data, batch_label, params) # We update the params
  }

}

# Test the accuracy of the neural network on the test data and labels
test = function(train_data, train_label, test_data, test_label, params) {
  pred_label <- forwardPropogate(test_data, params)[[2]]
  pred_label <- apply(pred_label, 1, which.max)
  acc <- mean(pred_label == test_label)
  return(acc)
}

# Test the accuracy of the neural network on the test data

```

```
acc <- test(train_data, train_label, test_data, test_label, params)
cat("Accuracy:", acc, "\n")
```

```
## Accuracy: 0.1885965
```

From a mathematical perspective, the gradient computation can be understood using the chain rule of calculus and the definition of the loss function.

Let's denote the loss function as L , which is a function of the predicted output z_2 and the true labels $label$. The goal is to compute the gradient of the loss function with respect to the weights of the second layer W_2 , denoted as

$$\frac{\partial L}{\partial W_2}$$

Using the chain rule, we can write the gradient of the loss function with respect to W_2 as:

$$\frac{\partial L}{\partial W_2} = \frac{\partial L}{\partial z_2} \cdot \frac{\partial z_2}{\partial W_2}$$

Here,

$$\frac{\partial L}{\partial z_2}$$

represents the gradient of the loss function with respect to the predicted output z_2 , and

$$\frac{\partial z_2}{\partial W_2}$$

represents the gradient of the predicted output z_2 with respect to the weights W_2 .

Now, let's assume the loss function is a mean squared error (MSE) between the predicted output z_2 and the true labels $label$, which is a common choice in neural networks:

$$L = \frac{1}{2} (label - z_2)^2$$

Using the chain rule, we can compute the gradient of the loss function with respect to z_2 as:

$$\frac{\partial L}{\partial z_2} = -(label - z_2)$$

This is because derivative of the squared error with respect to z_2 is

$$-2 \cdot (label - z_2)$$

, and we divide by 2 to get the average error.

Now, we need to compute the gradient of the predicted output z_2 with respect to the weights W_2 . Assuming the neural network has a linear activation function in the second layer, we can write z_2 as:

$$z_2 = W_2 \cdot z_1 + b_2$$

where z_1 is the output of the first layer, and b_2 is the bias term in the second layer.

Using the product rule, we can compute the gradient of z_2 with respect to W_2 as:

$$\frac{\partial z_2}{\partial W_2} = z_1$$

This is because the derivative of

$$W_2 \cdot z_1$$

with respect to W_2 is z_1 , and the derivative of b_2 with respect to W_2 is 0.

Now, we can plug in the expressions for

$$\frac{\partial L}{\partial z_2}$$

and

$$\frac{\partial z_2}{\partial W_2}$$

into the chain rule formula:

$$\frac{\partial L}{\partial W_2} = \frac{\partial L}{\partial z_2} \cdot \frac{\partial z_2}{\partial W_2} = -(\text{label} - z_2) \cdot z_1$$

This is exactly what the line of code `dw2 <- t(z1) %% db2_temp` is computing! The transpose of z_1 is used to ensure the correct matrix multiplication, and `db2_temp` represents the error term `label - z2`.

By computing the gradient of the loss function with respect to the weights W_2 in this way, we can update the weights using an optimization algorithm such as stochastic gradient descent (SGD) to minimize the loss function.

The chain rule is a fundamental concept in calculus that allows us to compute the derivative of a composite function. In this case, we have a composite function:

$$L = L(z_2)$$

where L is the loss function, and z_2 is the output of the second layer.

The chain rule states that if we have a composite function $f(g(x))$, where f and g are both functions of x , then the derivative of f with respect to x is given by:

$$\frac{\partial f}{\partial x} = \frac{\partial f}{\partial g} \cdot \frac{\partial g}{\partial x}$$

In our case, we can apply the chain rule by identifying the inner function g as z_2 , and the outer function f as L . Specifically, we can write:

$$L = L(z_2(W_2, z_1, b_2))$$

where z_2 is a function of W_2 , z_1 , and b_2 .

Now, we want to compute the derivative of L with respect to W_2 . Using the chain rule, we can write:

$$\frac{\partial L}{\partial W_2} = \frac{\partial L}{\partial z_2} \cdot \frac{\partial z_2}{\partial W_2}$$

Here, $\frac{\partial L}{\partial z_2}$ is the derivative of the loss function L with respect to the output z_2 , and $\frac{\partial z_2}{\partial W_2}$ is the derivative of the output z_2 with respect to the weights W_2 .

The key insight is that we can separate the computation of the derivative parts:

1. Compute the derivative of the loss function L with respect to the output z_2 . This is a simple computation, since L is a function of z_2 only.

$$\frac{L}{\partial z_2} = -(y - z_2)$$

where y is the true label.

2. Compute the derivative of the output z_2 with respect to the weights W_2 . This is also a simple computation, since z_2 is a function of W_2 , z_1 , and b_2 .

$$\frac{\partial z_2}{\partial W_2} = z_1$$

By multiplying these two derivatives together, we get the final result:

$$\frac{\partial L}{\partial W_2} = \frac{\partial L}{\partial z_2} \cdot \frac{\partial z_2}{\partial W_2} = -(y - z_2) \cdot z_1$$

This is exactly what the line of code `dw2 <- t(z1) %*% db2_temp` is computing! The transpose of z_1 is used to ensure the correct matrix multiplication, and `db2_temp` represents the error term $y - z_2$.

Here is the rewritten explanation with LaTeX expressions:

When computing the derivative of the loss function with respect to the weights W_2 , we can ignore the inputs z_1 and b_2 because they are not directly affected by the weights W_2 .

To see why, let's revisit the expression for the output z_2 :

$$z_2 = W_2 \cdot z_1 + b_2$$

When computing the derivative of z_2 with respect to W_2 , we only care about how W_2 affects the output z_2 . The inputs z_1 and b_2 are treated as constants, because they are not functions of W_2 .

In other words, when we compute the derivative of z_2 with respect to W_2 , we are asking: "How does W_2 change the output z_2 , assuming z_1 and b_2 are fixed?"

Using the product rule, we can compute the derivative of z_2 with respect to W_2 as:

$$\frac{\partial z_2}{\partial W_2} = z_1$$

Here, we've ignored the contribution of b_2 because it's a constant term that doesn't depend on W_2 . Similarly, we've treated z_1 as a constant, because it's an input that's not directly affected by W_2 .

By ignoring z_1 and b_2 in this way, we're effectively assuming that they are fixed inputs that don't change when we update the weights W_2 . This is a reasonable assumption, because we're only interested in how the weights W_2 affect the output z_2 , not how the inputs z_1 and b_2 affect the output.

Of course, this assumption is only valid if the inputs z_1 and b_2 are indeed fixed and don't depend on the weights W_2 . In general, if the inputs do depend on the weights, we would need to take that into account when computing the derivative.

But in this specific case, ignoring z_1 and b_2 allows us to focus on the direct effect of W_2 on the output z_2 , which is exactly what we need to compute the gradient of the loss function with respect to the weights.